


Alice, Bob, and the problem of distributing
software – a DevOps tale

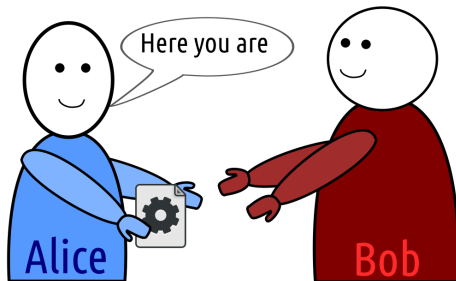
The core problem

- ▶ *Alice made a software tool*
- ▶ *Bob wants to run Alice's software*
- ▶ How can Alice make her software available to Bob?
- ▶ How can Bob run Alice's software?




Solving the problem?

- ▶ Alice gives Bob an executable
- ▶ Bob runs the executable



The real problems begin... (1/2)

Problem solved?



- ▶ Is Alice's executable suited for Bob's OS and architecture?
 - ▶ No? Alice must provide multiple executables for different systems
 - ▶ No? Bob must use *virtualisation* to emulate another OS/architecture.
 - ▶ Not compiled but interpreted? Bob must have the necessary interpreter (e.g. Python, JS, Java)

Back to the origin: Open Source

- ▶ Give Bob the source code and let him build the program himself?
- ▶ Ideally via a source version-controlled repository (e.g. GitHub, GitLab, BitBucket)
- ▶ Requires a higher level of technical expertise from Bob



The real problems continue: Dependency Hell (2/2)

- ▶ Does Alice's program make use of any other software (libraries) Bob also needs?
 - ▶ Alice must either provide or make explicit all her dependencies
 - ▶ Bob must ensure he has all the necessary dependencies before the program can run
- ▶ Alice's program may want to interact with other applications on the system?
 - ▶ Bob must ensure they're installed and set up properly
 - ▶ Related: Static linking vs Dynamic linking
- ▶ Recursive problem, conflicting requirements: **Dependency Hell**



Traditional solution: packaging and distributions

- ▶ Distributions host packages for common software in a *package repository*
- ▶ Packages in a distribution are carefully tuned to interoperate with one-another (ABI/API compatibility etc)
- ▶ A package manager handles packages and all their dependencies (apt, yum, apk, pacman, brew)
- ▶ Alice can now build a package (deb, rpm, apk) and add it to a repository from which Bob can install it.



Problem solved?

- ▶ No, there are many different distributions -> many package repositories
- ▶ Maintaining packages takes time and effort (for Alice)
- ▶ Packages have to be kept up to date (for both Bob and Alice)


Language ecosystems

- ▶ Language-ecosystems provide their own package repositories:
 - ▶ *Python* - Python Package Index - `pip`
 - ▶ *NodeJS* - `npm`
 - ▶ *Rust* - `crates.io` - `cargo`
 - ▶ *Java* - Maven Central - `maven`
 - ▶ *perl* - CPAN
 - ▶ *R* - CRAN
 - ▶ *Ruby* - Rubygems - `gem`
- ▶ If Alice's software fits into one of the ecosystems well, she should provide a package there

Software complexity: layers upon layers


Problem solved?

- ▶ No, Alice's software application may be an integration of an interconnected set of diverse software components:
 - ▶ Various languages: doesn't fit a single language ecosystem
 - ▶ Various audiences: doesn't fit a single distribution
- ▶ Though maybe captured in multiple more traditional packages, **configuring** the integration is often not trivial



Software-as-a-Service (SaaS)

- ▶ Instead of giving Bob the actual program, Alice makes her software available **as a service**.
- ▶ Bob can simply access it as a **web application** through his browser (or programmatically interact with it as a **web service**)
- ▶ Alice now gives her program to Charlie instead of Bob, to **deploy** on a server, Bob accesses Charlie's server.

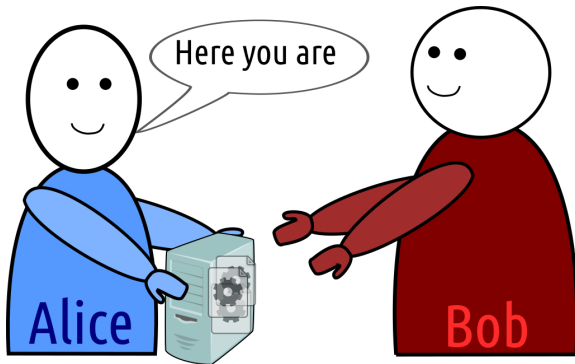


Problem solved?

- ▶ The burden shifts from Alice and Bob to Alice and Charlie (who has more technical expertise)
- ▶ **Service-as-a-Software Substitute**: having access to a service is a convenience but is **NOT** a substitute for having the actual software.
 - ▶ Privacy, data ownership and trust concerns
 - ▶ Low-level interoperability hindered: Increased latency, requires network connection
 - ▶ Business model vs technical solution
 - ▶ Can be a loophole to not provide the source code anymore (GPL vs AGPL)
 - ▶ Not suited for everybody

Virtualisation


- ▶ Bob: “*your program doesn't run*”
- ▶ Alice: “*But it works on my machine*”
- ▶ **Solution:** Alice just gives Bob a copy of her machine and Bob's machine *emulates* Alice's machine: a **virtual machine**





- ▶ Takes us back to scenario 1, *“Alice gives Bob an executable”*, but in the broadest possible sense of executable.
- ▶ Bob now only requires a **hypervisor** to run the VM.
- ▶ **Advantages:** Bridges OS differences, isolation from host (security)
- ▶ **Problem solved?** No
 - ▶ Performance penalty
 - ▶ Duplication
 - ▶ Resource overhead
 - ▶ Isolation from host (integration)

Containerisation

- ▶ Alice again gives Bob a copy of her machine: a **container**.



- ▶ Share the (Linux) kernel; isolate the rest, no virtualisation needed
- ▶ Takes advantage of facilities in the (Linux) kernel
- ▶ Significantly reduces resource overhead & performance penalty
- ▶ *Caveat:* container engines may resort to virtualisation anyway on non-Linux platforms (e.g. Docker on Windows or macOS)



Container platforms and paradigms (1/3)

Docker:

- ▶ Application containers, lightweight containers serving ideally single applications
- ▶ Non-persistent storage
- ▶ Layered images
- ▶ Containers are stateless (spun upon anew), all persistent data on externally mounted volumes
- ▶ Docker Hub as Container Registry (image store)
- ▶ Great for applications containers, deployment of services in cloud infrastructure
- ▶ Most known and widely used
- ▶ Less suitable for High Performance Clusters (security)

Container platforms and paradigms (2/3)

LXC:

- ▶ System containers: acts like a lightweight VM, more traditional environment
- ▶ Single-layer image
- ▶ File-system neutral
- ▶ Persistent storage (typically)
- ▶ Fat containers that may serve multiple applications, have a full *init system*.

Container platforms and paradigms (3/3)

Singularity:

- ▶ Specifically designed for High Performance Clusters and multi-user environments
- ▶ Can use docker images
- ▶ No elevated permissions required (security)


Provisioning a container (or VM)

- ▶ *"Alice gives Bob a copy of her machine"*
- ▶ Alice builds a (container/VM) image
- ▶ using an automated recipe for provisioning (e.g. Dockerfile, shell script, Ansible)
- ▶ Alice publishes the container image in a registry
- ▶ Bob obtains the container image from the registry



Container Orchestration

- ▶ Alice decided on Software-as-a-Service
- ▶ Alice gives Charlie a copy of her machine: a **container**, to offer as a service for Bob and others.



Problem solved?

- ▶ Alice's service may comprise multiple containers that need to interact.
 - ▶ L'histoire se répète?
- ▶ Alice's service may be so popular that running it on a single system is not sufficient (scalability)
- ▶ What if Alice's program fails and the service goes down?
- ▶ **Distributed Computing:** containers can easily be deployed on multiple systems
 - ▶ Schedule when a container runs and where it runs (load balancing etc)
 - ▶ Spin up multiple containers at once (e.g. docker compose)
 - ▶ Spin them up over multiple machines (e.g. docker swarm, kubernetes)
 - ▶ Restart a container when it fails (kubernetes)
 - ▶ Abstract over the hardware (kubernetes)

Recap

- ▶ Distribution
- ▶ Packaging
- ▶ Software as a Service
- ▶ Deployment
- ▶ Virtualisation
- ▶ Containerisation
- ▶ Container Orchestration

Recommendations for the CLARIAH infrastructure

No single solutions fits all, be aware of all the different layers and audiences.

For all individual software components:

- ▶ Try to make software components minimal components that are reusable
- ▶ Publish the source code in **public** Version Controlled Source Repositories (e.g. GitHub)
- ▶ Include software metadata with the source code (as codemeta or language-specific)
 - ▶ All direct dependencies must be made explicit
- ▶ Release software versions periodically when deemed stable
- ▶ Package and distribute all the software through proper channels:
 - ▶ Language-ecosystem-specific repositories (PyPi, CRAN, Maven etc)
 - ▶ or for specific distributions if appropriate

For software services (SaaS):

(Offering a service is perfectly valid, and one of the aims of CLARIAH, but is **not** a substitute for offering the actual software that the user can run himself/herself)!

For software services (SaaS):

- ▶ Software services should be delivered to service providers in **containers**:
 - ▶ Docker is the most industry-standard solution:
 - ▶ Software **developers** should provide a Dockerfile (in public version-controlled source repository) allowing anyone to build a container/ multiple containers to deliver the service
 - ▶ Application containers (hosting a single application) should be preferred
 - ▶ CLARIAH should set up a container registry for use by all participants
 - ▶ All service endpoints that a container exposes should be clearly documented
 - ▶ When multiple interacting containers are delivered, an initial Compose file or Kubernetes configuration should be provided by the software **developers** to orchestrate them.
 - ▶ These should be stored in code in version control (may be private).
 - ▶ CLARIAH should draft clear infrastructure requirements: the NDE example is excellent to copy and follow, which is in turn partially based on the Twelve-Factors.

References

- ▶ NDE Infrastructure requirements: <https://github.com/netwerk-digitaal-erfgoed/requirements-infrastructure>
- ▶ Twelve Factors: <https://12factor.net/>
- ▶ Clipart: openclipart.org (heavily modified by me without any real artistic talent)
- ▶ XKCD (comics)